



BUILD GUIDE B · GENERAL BUSINESS DATA

Find and pull out

Messy documents in. The specific answer out. A fictional 14-person Philadelphia design and web studio. No financial vocabulary anywhere. Everything you need is in this one document.

Event

Finance of the Future

Date

Wednesday 29 July 2026

Data set

Fairmount Studio

01 / DIAGNOSIS

Is this your problem?

Messy documents go in. The specific answer comes out. The answer already exists somewhere. Finding it is the entire job, and finding it is what takes the afternoon.

- ☐ The information you need is already written down. It is just buried.
- ☐ It is spread across files, emails, PDFs, a shared drive, or a portal.
- ☐ You are looking for the same handful of things every time, in different documents.
- ☐ Being wrong matters. This is not a "roughly right is fine" task.

Typical shapes: finding the right file in SharePoint or Drive, pulling key terms out of a stack of agreements, checking data quality across exports, building a comparison table from eleven statements, extracting spend from a year of email receipts, reading a data room.

READ THIS BEFORE YOU BUILD ANYTHING

This is the cluster where output **looks right and is wrong**. A weekly report that drifts is obvious at a glance. An extraction table with one silently invented cell looks identical to a correct one. Everything in this guide is built around that single problem. Do not skip Section 06.

02 / START HERE

Two boxes. That is the whole thing.

Every tool has two places to type. Getting them the wrong way round is the most common mistake in the room, and it produces output that looks almost right and is useless tomorrow.

TOP BOX · PASTE ONCE, LEAVE IT THERE

The instruction

Google AI Studio: the panel on the right labelled *System instructions*. **Claude:** make a Project, then *Project instructions*. **ChatGPT:** make a Project, or a Custom GPT, then *Instructions*.

This is the thing you are building tonight. It stays put. You never retype it.



BOTTOM BOX · PASTE FRESH EVERY RUN

The data

The normal chat box at the bottom of the screen, where you would usually type a question.

This changes every time. Tonight it is workshop data. On Monday it is your real work.

Three steps. Everything you need is in this document.

01

Open your tool

aistudio.google.com is fastest if you have a Google account. claude.ai or chatgpt.com work identically.

02

Copy Section 03 of this guide into the TOP box

All of it, first line to last. Do not edit it. Do not fill anything in. It is already complete and ready to run.

03

Copy Section 04 of this guide into the BOTTOM box, then press run

Section 04 is an internal review memo from Fairmount Studio. One run is enough to start.

IF THE OUTPUT LOOKS STRANGE, CHECK WHICH BOX YOU USED

Instruction in the bottom box means you asked a one-off question and built nothing reusable. Data in the top box means every future run will be contaminated with tonight's numbers. Both produce something that reads fine and is wrong. This one mistake accounts for most of the confusion in any workshop like this, and it takes ten seconds to check.

YOU DO NOT NEED TO OPEN ANYTHING ELSE

The instruction, the data and the answer key are all in this document. Run it as written first, even though it is about a firm that is not yours. Seeing the pattern work end to end takes four minutes and makes every decision after it obvious. Section 07 turns it into your own.

Copy this into the top box

Complete and ready to run. No brackets to fill, nothing to decide. Copy every line.

You extract reported operating metrics for a small studio from an internal review memo.

INPUT

I will paste an internal memo. It is prose, not a table. Figures are scattered through it. It may contain typos, repeated figures, and figures stated as of different dates.

OUTPUT

A single table. One row per metric. Exactly these columns:
| metric | value | unit | as-of date | source quote | check |

Extract only these metrics, in this order:

1. Projects delivered
2. Average project value
3. Billable utilisation
4. Revenue
5. Headcount
6. Average days to invoice payment
7. Client retention rate
8. Cost per new client acquired

- "source quote" is the exact sentence the value came from. Verbatim.
- "as-of date" comes from the memo. If not stated, write "not stated".
- "check" is empty unless you are unsure.

After the table, output a section called "Not found" listing every requested metric that does not appear, and a section called "Conflicts" listing anything that appears more than once with different values.

RULES

- Every value must come from the memo. Do not use anything you know about this firm from outside the memo.
- Copy figures exactly as written. Do not convert \$1.28 million into 1,280,000. Do not round.
- Output the table only. No commentary.

REFUSALS

- If a metric is not in the memo, write "not found". Never infer, never estimate, never derive it from another figure.
- If the same metric appears twice with different values, output both rows and list it under Conflicts. Do not pick one.
- If two values differ because they are as of different dates, that is not a conflict. Record both with their dates.

Section 08 has the same instruction with the specifics stripped out and brackets in their place, for when you point this at your own work.

Copy this into the bottom box

An internal memo from Fairmount Studio, written the way real internal memos are written. The numbers are scattered through prose, some are repeated, and one of them is wrong.

FAIRMOUNT STUDIO
INTERNAL MEMORANDUM. NOT FOR CIRCULATION OUTSIDE THE PARTNERS.
To: Partners
From: Studio manager
Date: 24 July 2026
Re: First half 2026 review, ahead of the planning session

Pulling together where we landed for the first half of the year, to 30 June 2026, so we are all working from the same numbers on Thursday. I have written this out rather than building a deck because half of it needs discussion rather than presentation.

On volume. We delivered 41 projects in the first half, which is up on the same period last year and is more than I expected given we were two people short until March. Average project value came out at \$14,200. That average is doing a lot of work and hides a real spread: four projects accounted for close to a third of the revenue, and the long tail of small jobs is where all the admin lives. Worth a conversation about a minimum engagement size.

On revenue. First half revenue was \$1.28 million. Client retention sat at 82 percent, which I am pleased with, though the three clients we lost were all in the small tail rather than the top end. Average days to invoice payment was 47 days across the half, which is worse than it should be and is the single thing most likely to cause us a cash problem in the autumn. Two clients account for most of that number.

On capacity. Headcount was 14 as of 30 June. Billable utilisation was 68 percent as of 30 June, and it has since risen to 71 percent as of 24 July as the summer projects kicked in. I want to flag both figures with their dates rather than quoting one, because they tell different stories and the July number is not yet a trend. We have also taken on two more people since the half closed, so headcount is now 16, which will pull utilisation down again in the short term before it helps.

Two things I could not resolve. First, my note from the March partners meeting has first half revenue at \$128,000, which is obviously not right alongside the \$1.28 million figure above, and I cannot work out which spreadsheet the error came from. Somebody should reconcile these before Thursday rather than in the room. Second, we still do not track what it costs us to win a new client. I have been asked for that number three times this year and I have never been able to produce it, because the marketing spend, the proposal hours and the referral relationships are recorded in three places and none of them reconcile. That is the actual finding in this memo.

Recommendation: fix the payment terms conversation before we touch anything else. Everything on this list is a lot easier with 30 days instead of 47.

05 / RUN IT AND CHECK

Did it actually work?


The memo contains three planted problems. A working instruction catches all three. Check the quotes, not the answers: pick two rows at random and search the memo for the quoted text. If it is not there verbatim, throw the run away and fix the instruction.

METRIC	CORRECT ANSWER	AS-OF	WHAT IT IS TESTING
Projects delivered	41	H1 2026, to 30 Jun 2026	Clean.
Average project value	\$14,200	H1 2026	Clean.
Billable utilisation	68% AND 71%. Two rows.	30 Jun 2026 / 24 Jul 2026	DATE TRAP. Both correct. Must NOT be flagged as a conflict. Both rows carry their date.
Revenue	\$1.28 million. The memo also says \$128,000.	H1 2026	TYPO CONFLICT. Both values must appear, listed under Conflicts. Picking one silently is the expensive failure.
Headcount	14	30 Jun 2026	Clean. The memo also gives 16 after two July hires. Different date, not a conflict.
Average days to invoice payment	47 days	H1 2026	Clean.
Client retention rate	82%	H1 2026	Clean.
Cost per new client acquired	NOT FOUND	n/a	ABSENT METRIC. Not in the memo at all. Any number here means the tool invents, and every other row is unverified.

THE THIRTY-SECOND AUDIT, FOREVER

Two rows at random. Search the source for the quoted text. If both quotes exist verbatim and support the value, ship it. Do this on every run for the first month. It costs thirty seconds and it is the entire quality system.

What you are building



The output is always a table, even when you only want one answer, because a table forces one row per item and one column per field, and that structure is what stops the model from writing a fluent paragraph that blurs three documents together.

Swap in your own work

You have a working version. Now point it at your own work. About twenty minutes, and the result is the thing you actually use on Monday.

MINUTES 0 TO 4

01

Name the columns before you name the documents

Write down every field you actually need, as a column heading. "Counterparty, effective date, fee basis, termination notice, governing law." Not "the key terms." If you cannot list the columns, you do not yet know what you are extracting, and the model will happily fill that gap with its own idea of what matters.

MINUTES 4 TO 8

02

Decide what proof looks like

For each row you want the model to quote the exact line it took the value from. This is called grounding and it is the single most effective quality control available to you. It converts "trust the output" into "spot check the output in thirty seconds", which is the difference between a tool you use and a tool you abandon.

MINUTES 8 TO 13

03

Paste the blank scaffold in Section 06, fill the brackets

Set the columns. Set the "not found" rule. Do not soften either.

MINUTES 13 TO 18

04

Run it on a document where you already know the answer

Critical. You cannot evaluate an extraction tool on a document you have not read. Use the source document in your pack, or one of your own that you know cold.

MINUTES 18 TO 25

05

Check the quotes, not the answers

Go to the source. Confirm the quoted lines exist, verbatim, and say what the table claims they say. Section 07.

The version with the brackets

You extract [WHAT] from [WHAT KIND OF DOCUMENT].

INPUT

I will paste or upload one or more [document type]. They are inconsistent in format and may contain errors, duplicates, or contradictions.

OUTPUT

A single table. One row per [thing]. Exactly these columns:

| [column] | [column] | [column] | as-of date | source quote | check |

- "source quote" is the exact text from the document that the values in that row came from. Verbatim. Include enough words to find it.
- "as-of date" is the date the figure applies to, taken from the document. If the document does not state one, write "not stated".
- "check" is empty unless you are unsure, in which case put a short reason.

After the table, output a section called "Not found" listing every requested field that does not appear in the source, and a section called "Conflicts" listing anything that appears more than once with different values.

RULES

- Every value must come from the document. Do not use anything you know about [subject] from outside the document.
- Do not convert units, round, or restate figures. Copy them in the form they appear.
- Do not combine values from two different documents into one row.
- Output the table only. No summary, no commentary, no preamble.

REFUSALS

- If a field is not present, write "not found". Never infer it, never estimate it, never derive it from a related figure.
- If the same field appears twice with different values, output both rows and list it under Conflicts. Do not pick one and do not average them.
- If two figures differ because they are as of different dates, that is not a conflict. Record both with their dates.

WHY "DO NOT USE ANYTHING YOU KNOW FROM OUTSIDE THE DOCUMENT"

Without that line, a model asked for a company's revenue will sometimes answer from its training data instead of from your file, and you will never know which one you got. The line is short, it costs nothing, and it is the difference between an extraction tool and a very confident encyclopedia.

When it goes wrong

WHAT YOU SEE

WHAT TO CHANGE

Source quotes are paraphrased, not verbatim.

Add: "The source quote must be copyable. If I search the document for that exact string I must find it." Then test by doing exactly that.

It fills gaps with reasonable-sounding values.

Move REFUSALS above OUTPUT and add: "Before producing the table, list every requested field you could not locate."

It merges two documents into one row.

Add a "source document" column. The moment it has to name the file, the merging stops.

It misses things that are clearly there.

The input is too long for one pass. Split it and run in batches, then combine. Do not try to fix this with wording.

It converts and rounds your figures.

"Copy figures exactly as written" plus one literal example of a figure you do not want converted.

It adds a helpful summary you did not ask for.

"Output the table only." Repeat the constraint at the very end of the instruction. Last-line instructions carry weight.

Different answers on identical input, run twice.

Your columns are ambiguous. If two careful humans would disagree about what belongs in a column, the model will too. Rewrite the column definition, not the prompt.

10 / FURTHER

Three rungs up, after tonight

01 Point it at the folder (connectors)

This is the cluster where connectors pay off most, because the pain is volume. With the Drive or SharePoint connector on, INPUT becomes "read every PDF in this folder" and the eleven-file download step disappears. Read access only. There is no reason an extraction tool should ever have write access.

02 Two-pass extraction (agentic workflow)

Pass one extracts. Pass two is a separate instruction whose only job is to audit pass one: "here is a table and here is the source document. For each row, confirm the quote exists verbatim and supports the value. Output only the rows that fail." A checker that never saw the extraction step catches things the extractor cannot. This is the highest-value upgrade in this guide.

03 Watch the folder (scheduled)

New file lands, extraction runs, row appends to a sheet, you get a message only when the check column is not empty. Silence means clean. This is the shape of every good monitoring workflow: it contacts you when something needs a human, and never otherwise.

MONDAY

Take the ten documents you were going to read this week. Run the extraction. Then read three of them properly and check the table against them. You are not testing whether it saves time. You are establishing your own error rate, which is the number you will need when someone at your firm asks whether this is safe.

Everything you need is in this document. The instruction (Section 03), the data (Section 04), the answer key (Section 05), and the blank version for your own work (Section 08). You do not need to open another file.

Fairmount Studio is not a real firm. All client names, staff, invoices, memos and inbox items are invented for this workshop. No real person appears anywhere in this data.

oneprompt. Philly's AI community for business. neilrayarao@oneprompt.live · oneprompt.live